




MACHINE LEARNING FOR FINANCE IN PYTHON

Engineering features

Nathan George
Data Science Professor



One problem with linear models

```
# add non-linear interaction term for a linear model  
SMAxRSI = amd_df['14-day SMA'] * amd_df['14-day RSI']
```

Some models that don't require manually creating interaction features:

Decision-tree-based models


- Random forests
- Gradient boosting

Others


- neural networks



Feature engineering

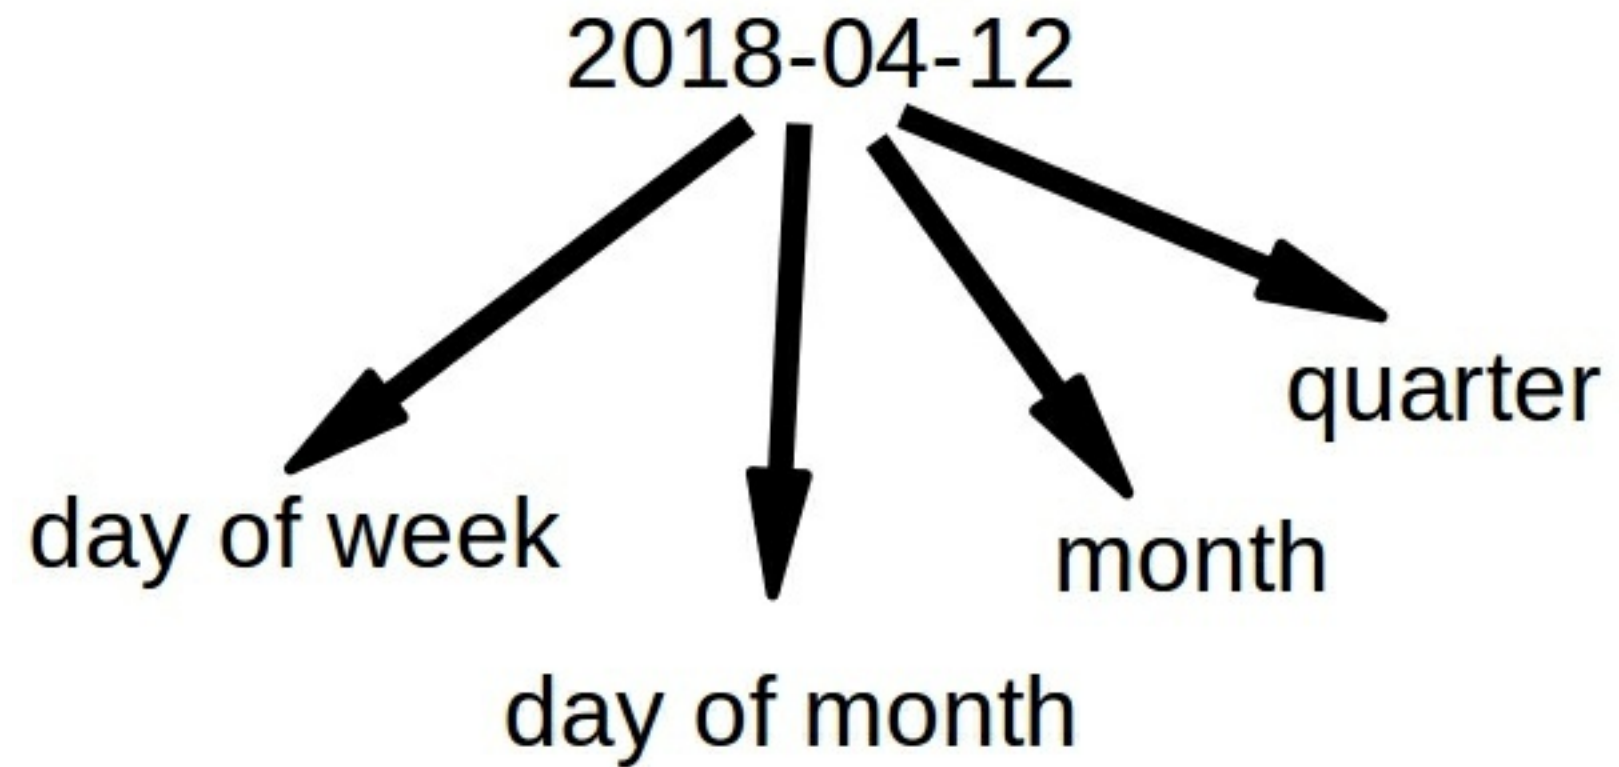


Volume





Datetime feature engineering





Extracting the day of week

```
print(amd_df.index.dayofweek)
```

```
Int64Index([2, 3, 4, 0, 1, 2, 3, 4, 0, 1,  
           ...  
           1, 2, 3, 4, 0, 1, 2, 3, 4, 0],  
           dtype='int64', name='Date', length=4807)
```





Dummies

```
days_of_week = pd.get_dummies(amd_df.index.dayofweek,  
                               prefix='weekday',  
                               drop_first=True)
```

```
print(days_of_week.head())
```

| | weekday_1 | weekday_2 | weekday_3 | weekday_4 |
|------------|-----------|-----------|-----------|-----------|
| Date | | | | |
| 2018-04-10 | 1 | 0 | 0 | 0 |
| 2018-04-11 | 0 | 1 | 0 | 0 |
| 2018-04-12 | 0 | 0 | 1 | 0 |
| 2018-04-13 | 0 | 0 | 0 | 1 |
| 2018-04-16 | 0 | 0 | 0 | 0 |





MACHINE LEARNING FOR FINANCE IN PYTHON

**Engineer some
features!**




MACHINE LEARNING FOR FINANCE IN PYTHON


Decision Trees

Nathan George
Data Science Professor

Decision trees




Decision trees






Decision tree splits

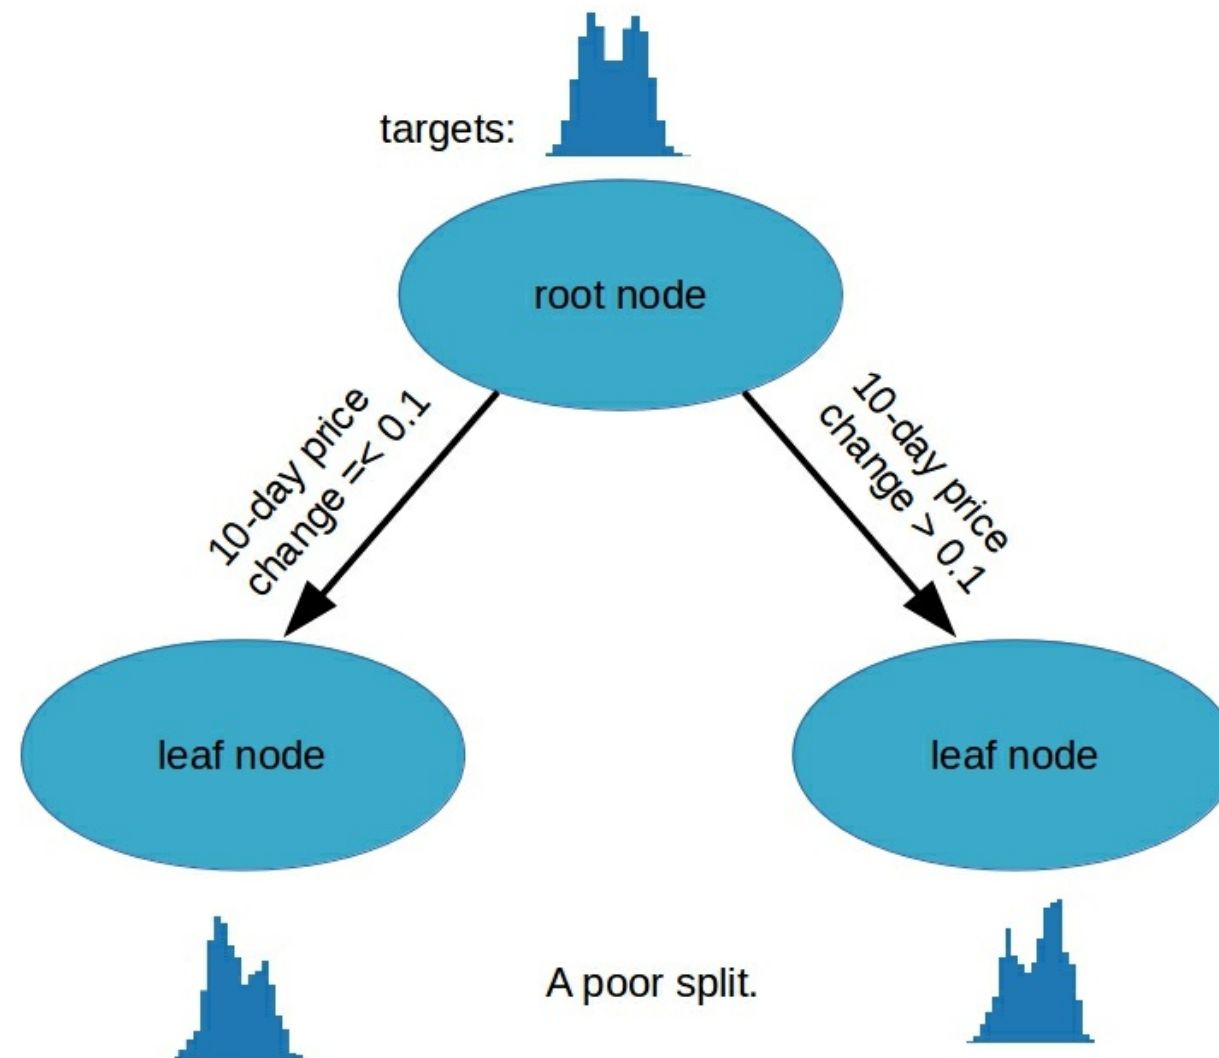




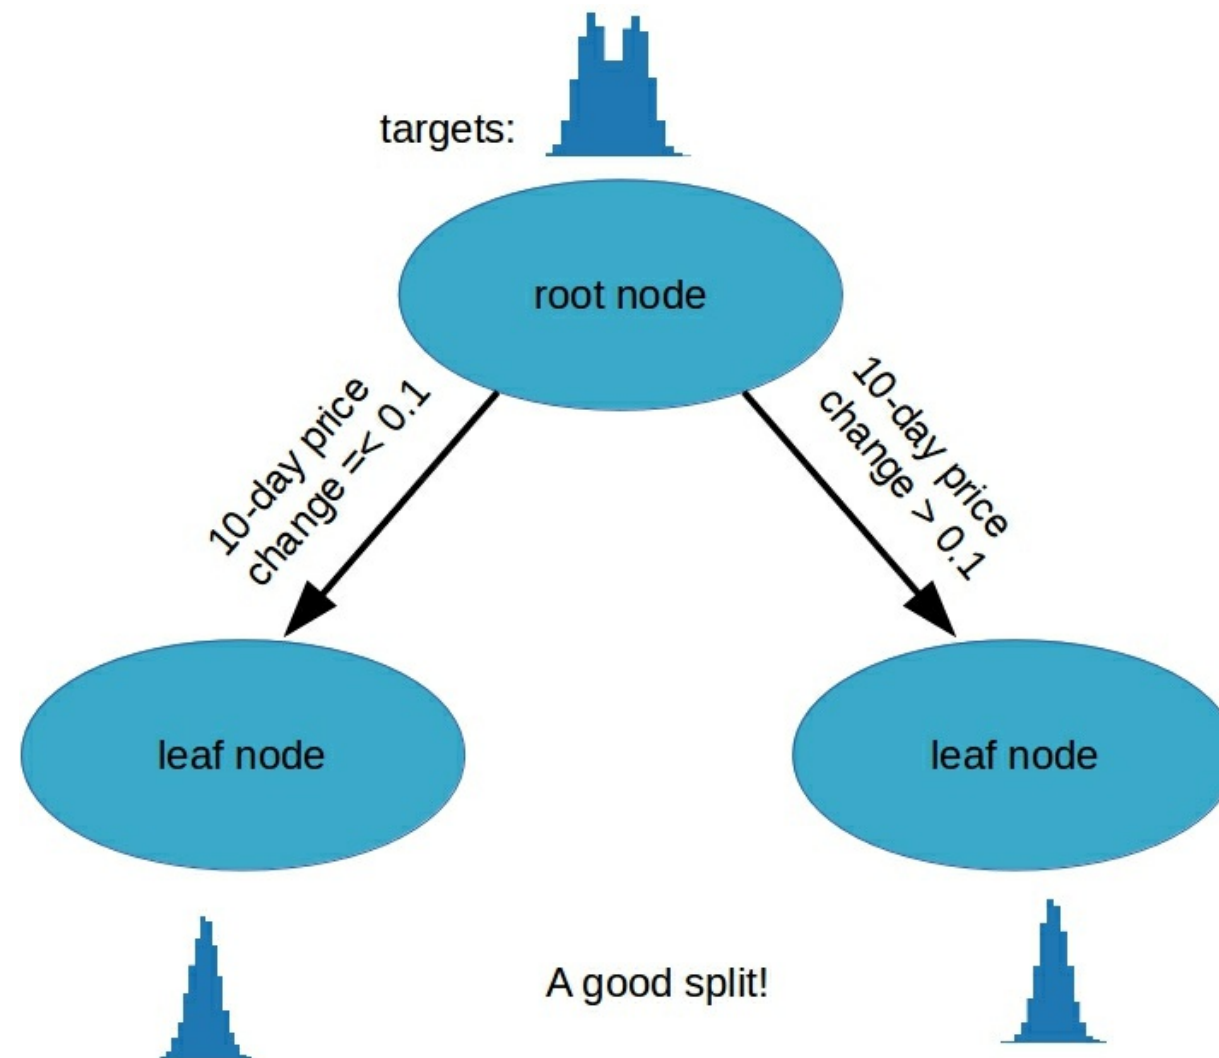
Decision tree splits



Bad tree

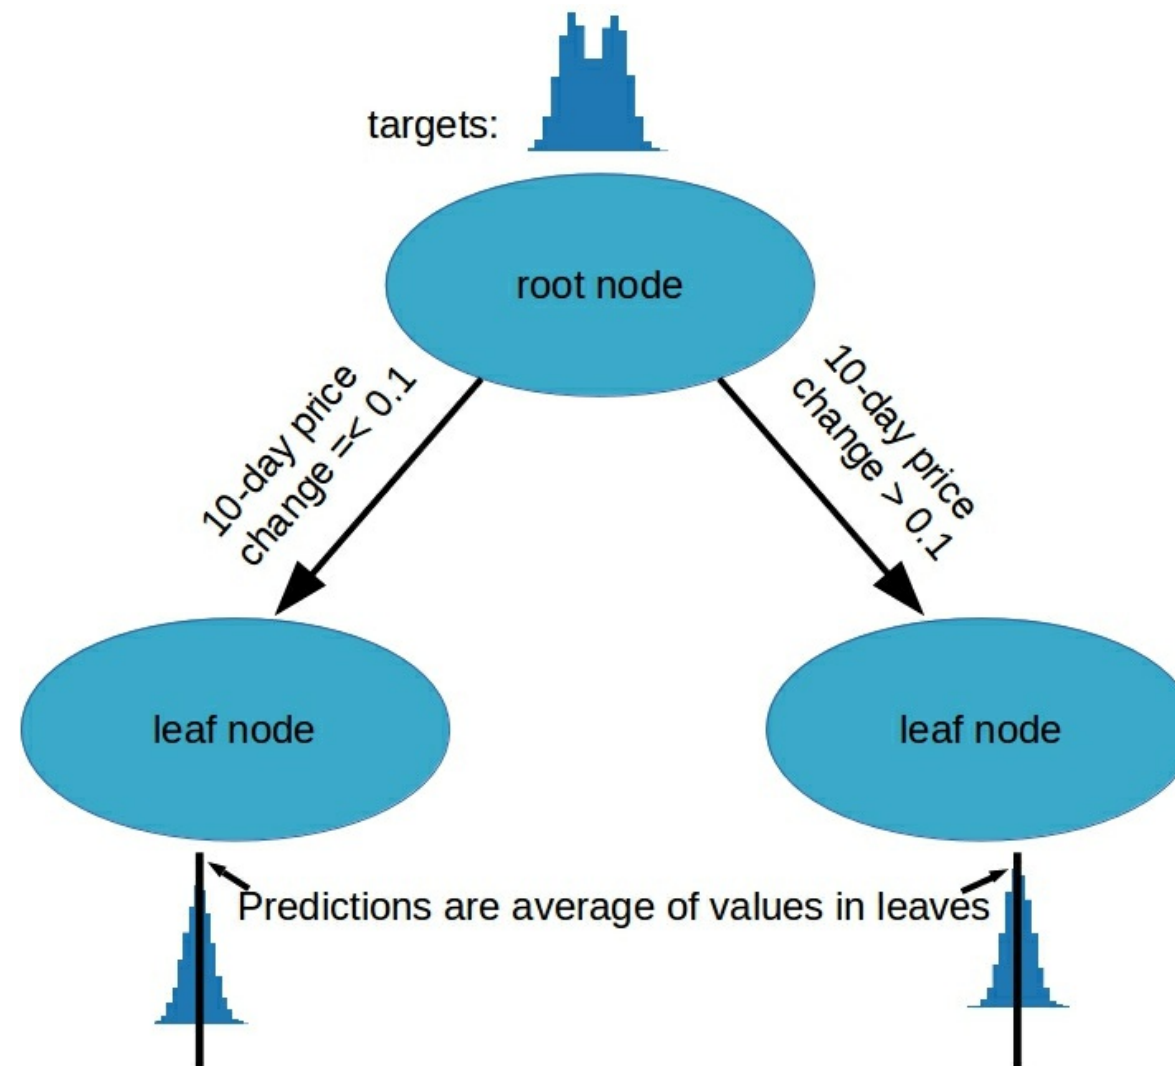


Good tree





Decision tree regression

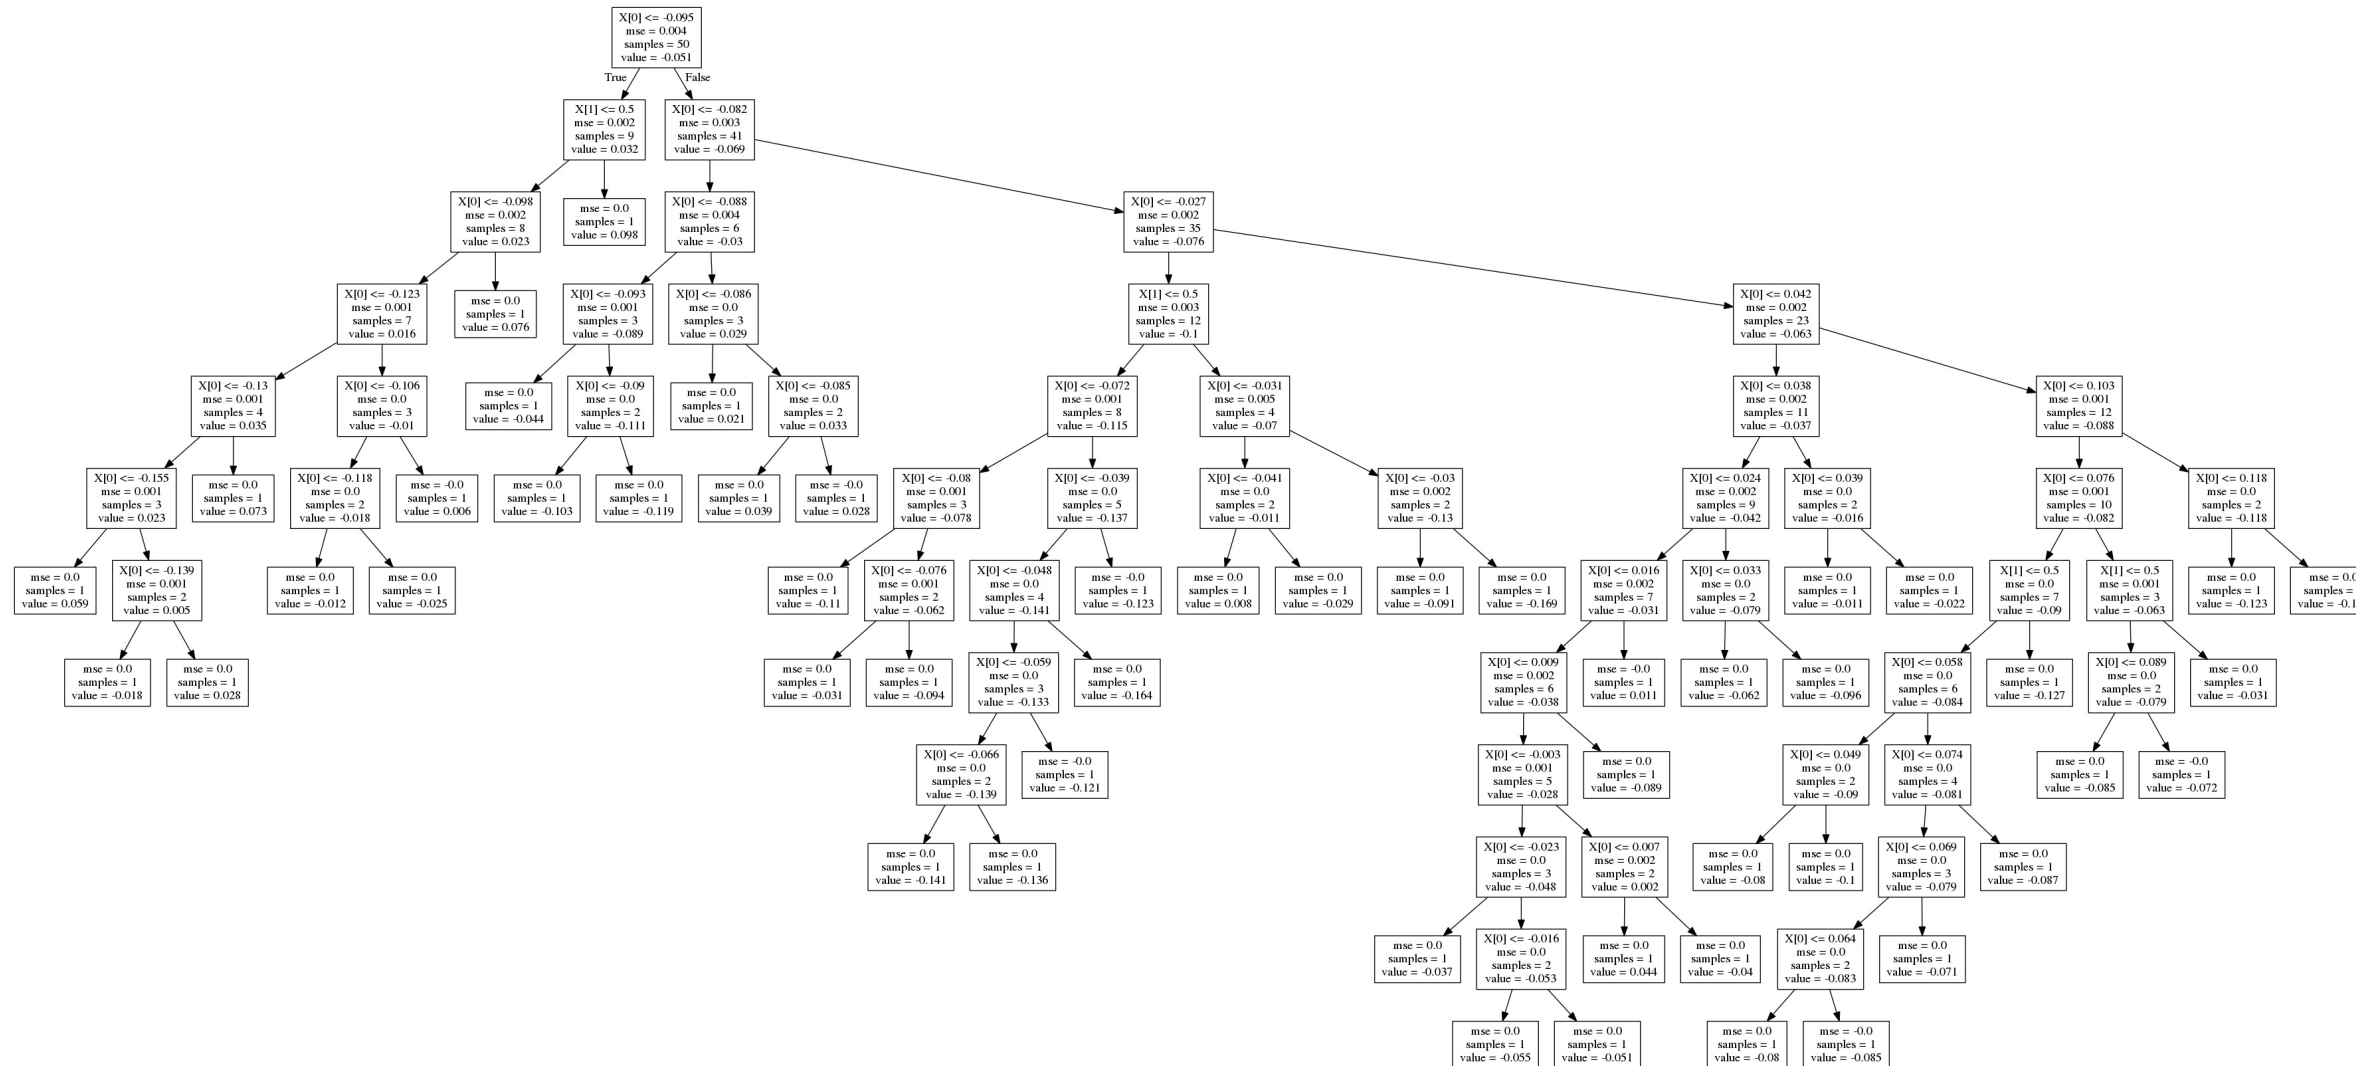




Regression trees

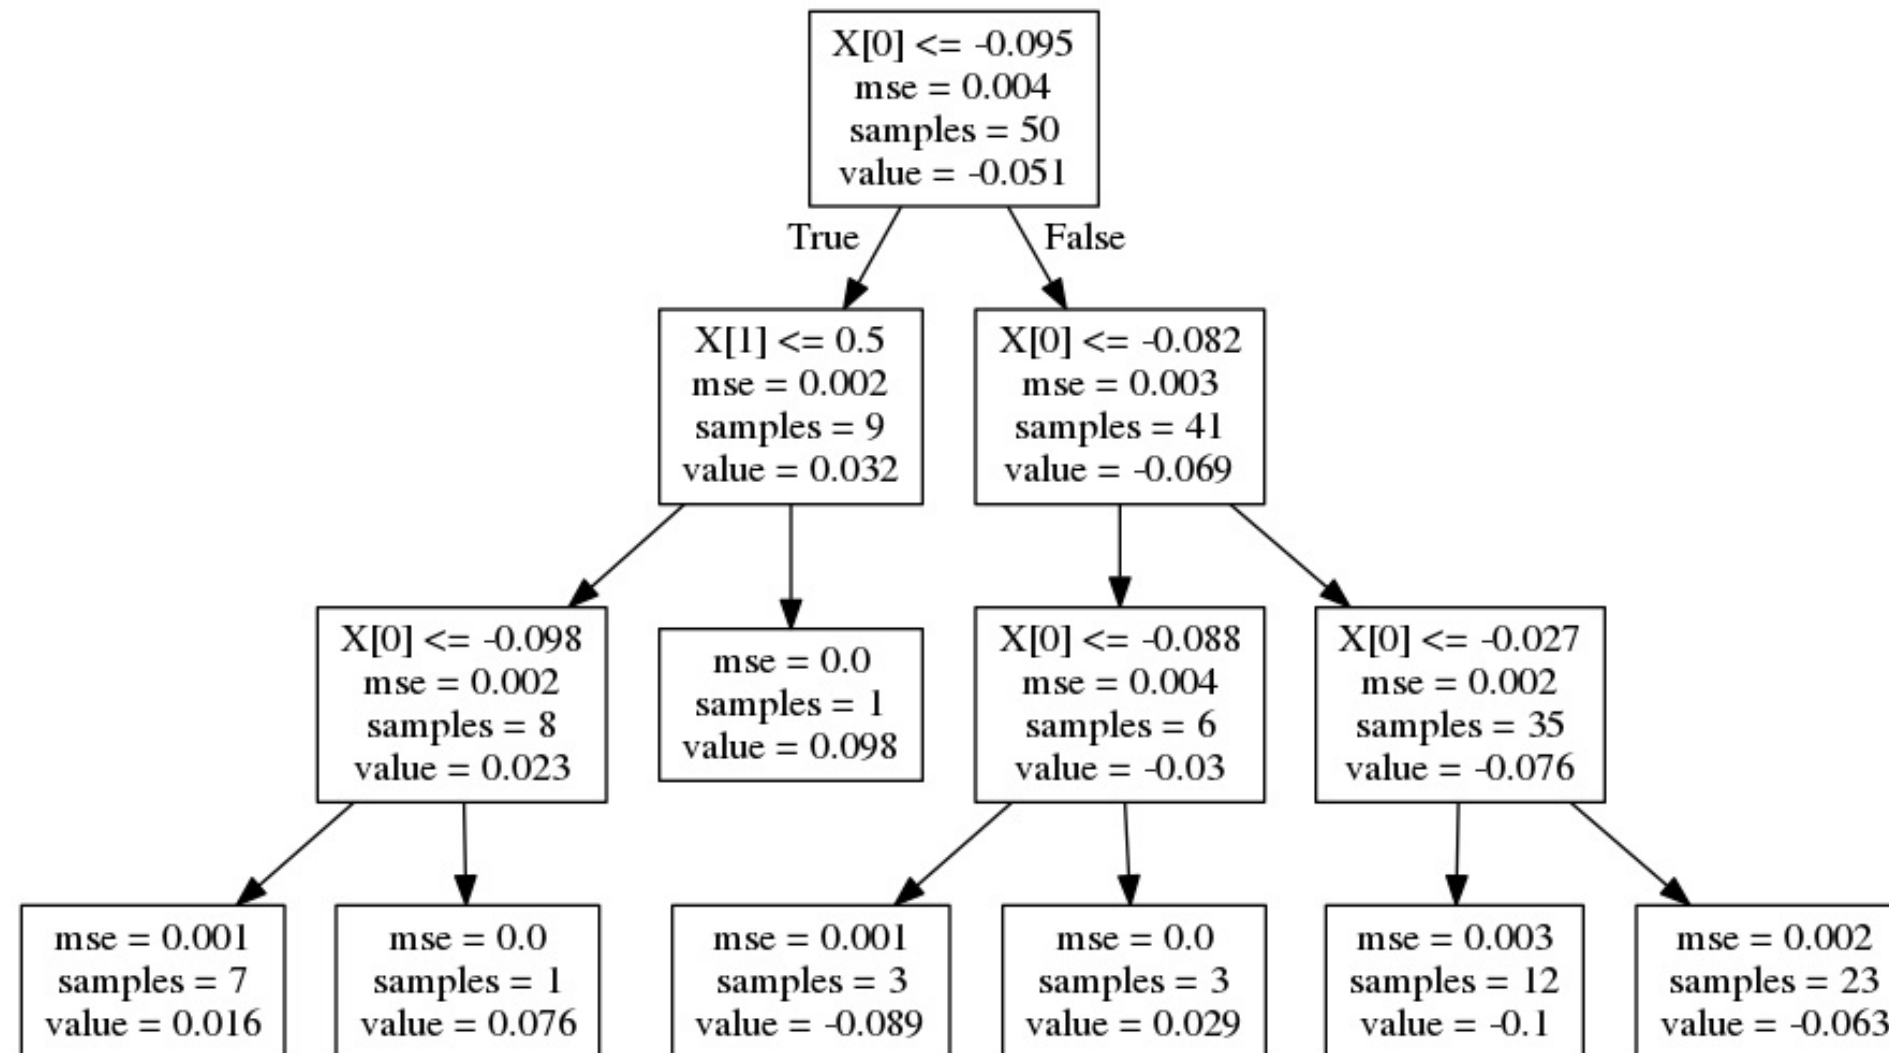
```
from sklearn.tree import DecisionTreeRegressor  
decision_tree = DecisionTreeRegressor(max_depth=5)  
decision_tree.fit(train_features, train_targets)
```

Decision tree hyperparameters





Max depth of 3






Evaluate model

```
print(decision_tree.score(train_features, train_targets))  
print(decision_tree.score(test_features, test_targets))
```

```
0.6662215501032416  
-0.08917300191734268
```

```
train_predictions = decision_tree.predict(train_features)  
test_predictions = decision_tree.predict(test_features)  
plt.scatter(train_predictions, train_targets, label='train')  
plt.scatter(test_predictions, test_targets, label='test')  
plt.legend()  
plt.show()
```





MACHINE LEARNING FOR FINANCE IN PYTHON


Grow some trees!




MACHINE LEARNING FOR FINANCE IN PYTHON

Random forests

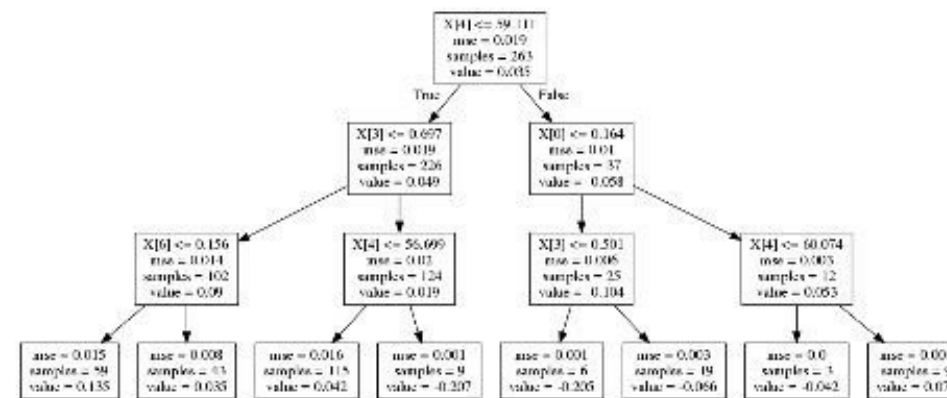
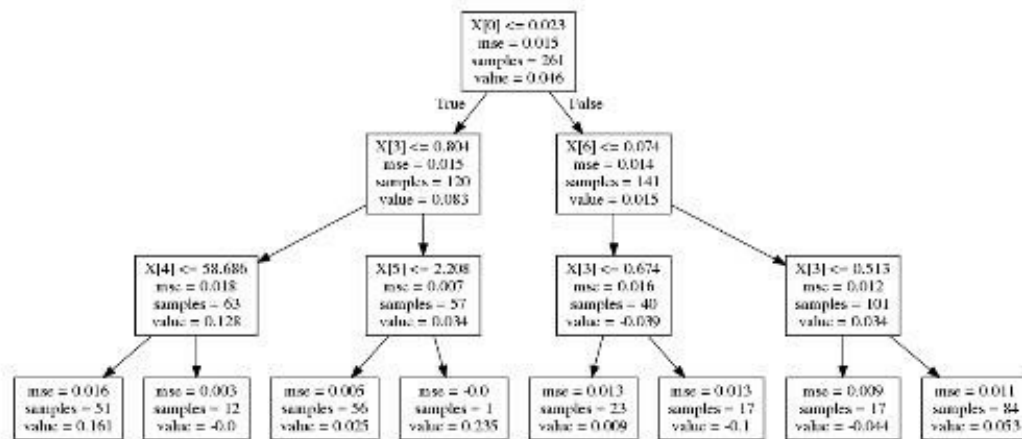


Nathan George
Data Science Professor








Random forests






Bootstrap aggregating (bagging)

training data



bootstrapped sample





Feature sampling

Random Forests

- A collection (ensemble) of decision trees
- Bootstrap aggregating (bagging)
- Sample of features at each split



sklearn implementation

```
from sklearn.ensemble import RandomForestRegressor

random_forest = RandomForestRegressor()
random_forest.fit(train_features, train_targets)
print(random_forest.score(train_features, train_targets))
```




Hyperparameters

```
random_forest = RandomForestRegressor(n_estimators=200,  
                                     max_depth=5,  
                                     max_features=4,  
                                     random_state=42)
```



ParameterGrid

```
from sklearn.model_selection import ParameterGrid

grid = {'n_estimators': [200], 'max_depth': [3, 5], 'max_features': [4, 8]}

from pprint import pprint

pprint(list(ParameterGrid(grid)))

[{'max_depth': 3, 'max_features': 4, 'n_estimators': 200},
 {'max_depth': 3, 'max_features': 8, 'n_estimators': 200},
 {'max_depth': 5, 'max_features': 4, 'n_estimators': 200},
 {'max_depth': 5, 'max_features': 8, 'n_estimators': 200}]
```

ParamaterGrid

```
test_scores = []

# loop through the parameter grid, set hyperparameters, save the scores
for g in ParameterGrid(grid):
    rfr.set_params(**g) # ** is "unpacking" the dictionary
    rfr.fit(train_features, train_targets)
    test_scores.append(rfr.score(test_features, test_targets))

# find best hyperparameters from the test score and print
best_idx = np.argmax(test_scores)
print(test_scores[best_idx])
print(ParameterGrid(grid)[best_idx])

0.05594252725411142
{'max_depth': 5, 'max_features': 8, 'n_estimators': 200}
```



MACHINE LEARNING FOR FINANCE IN PYTHON

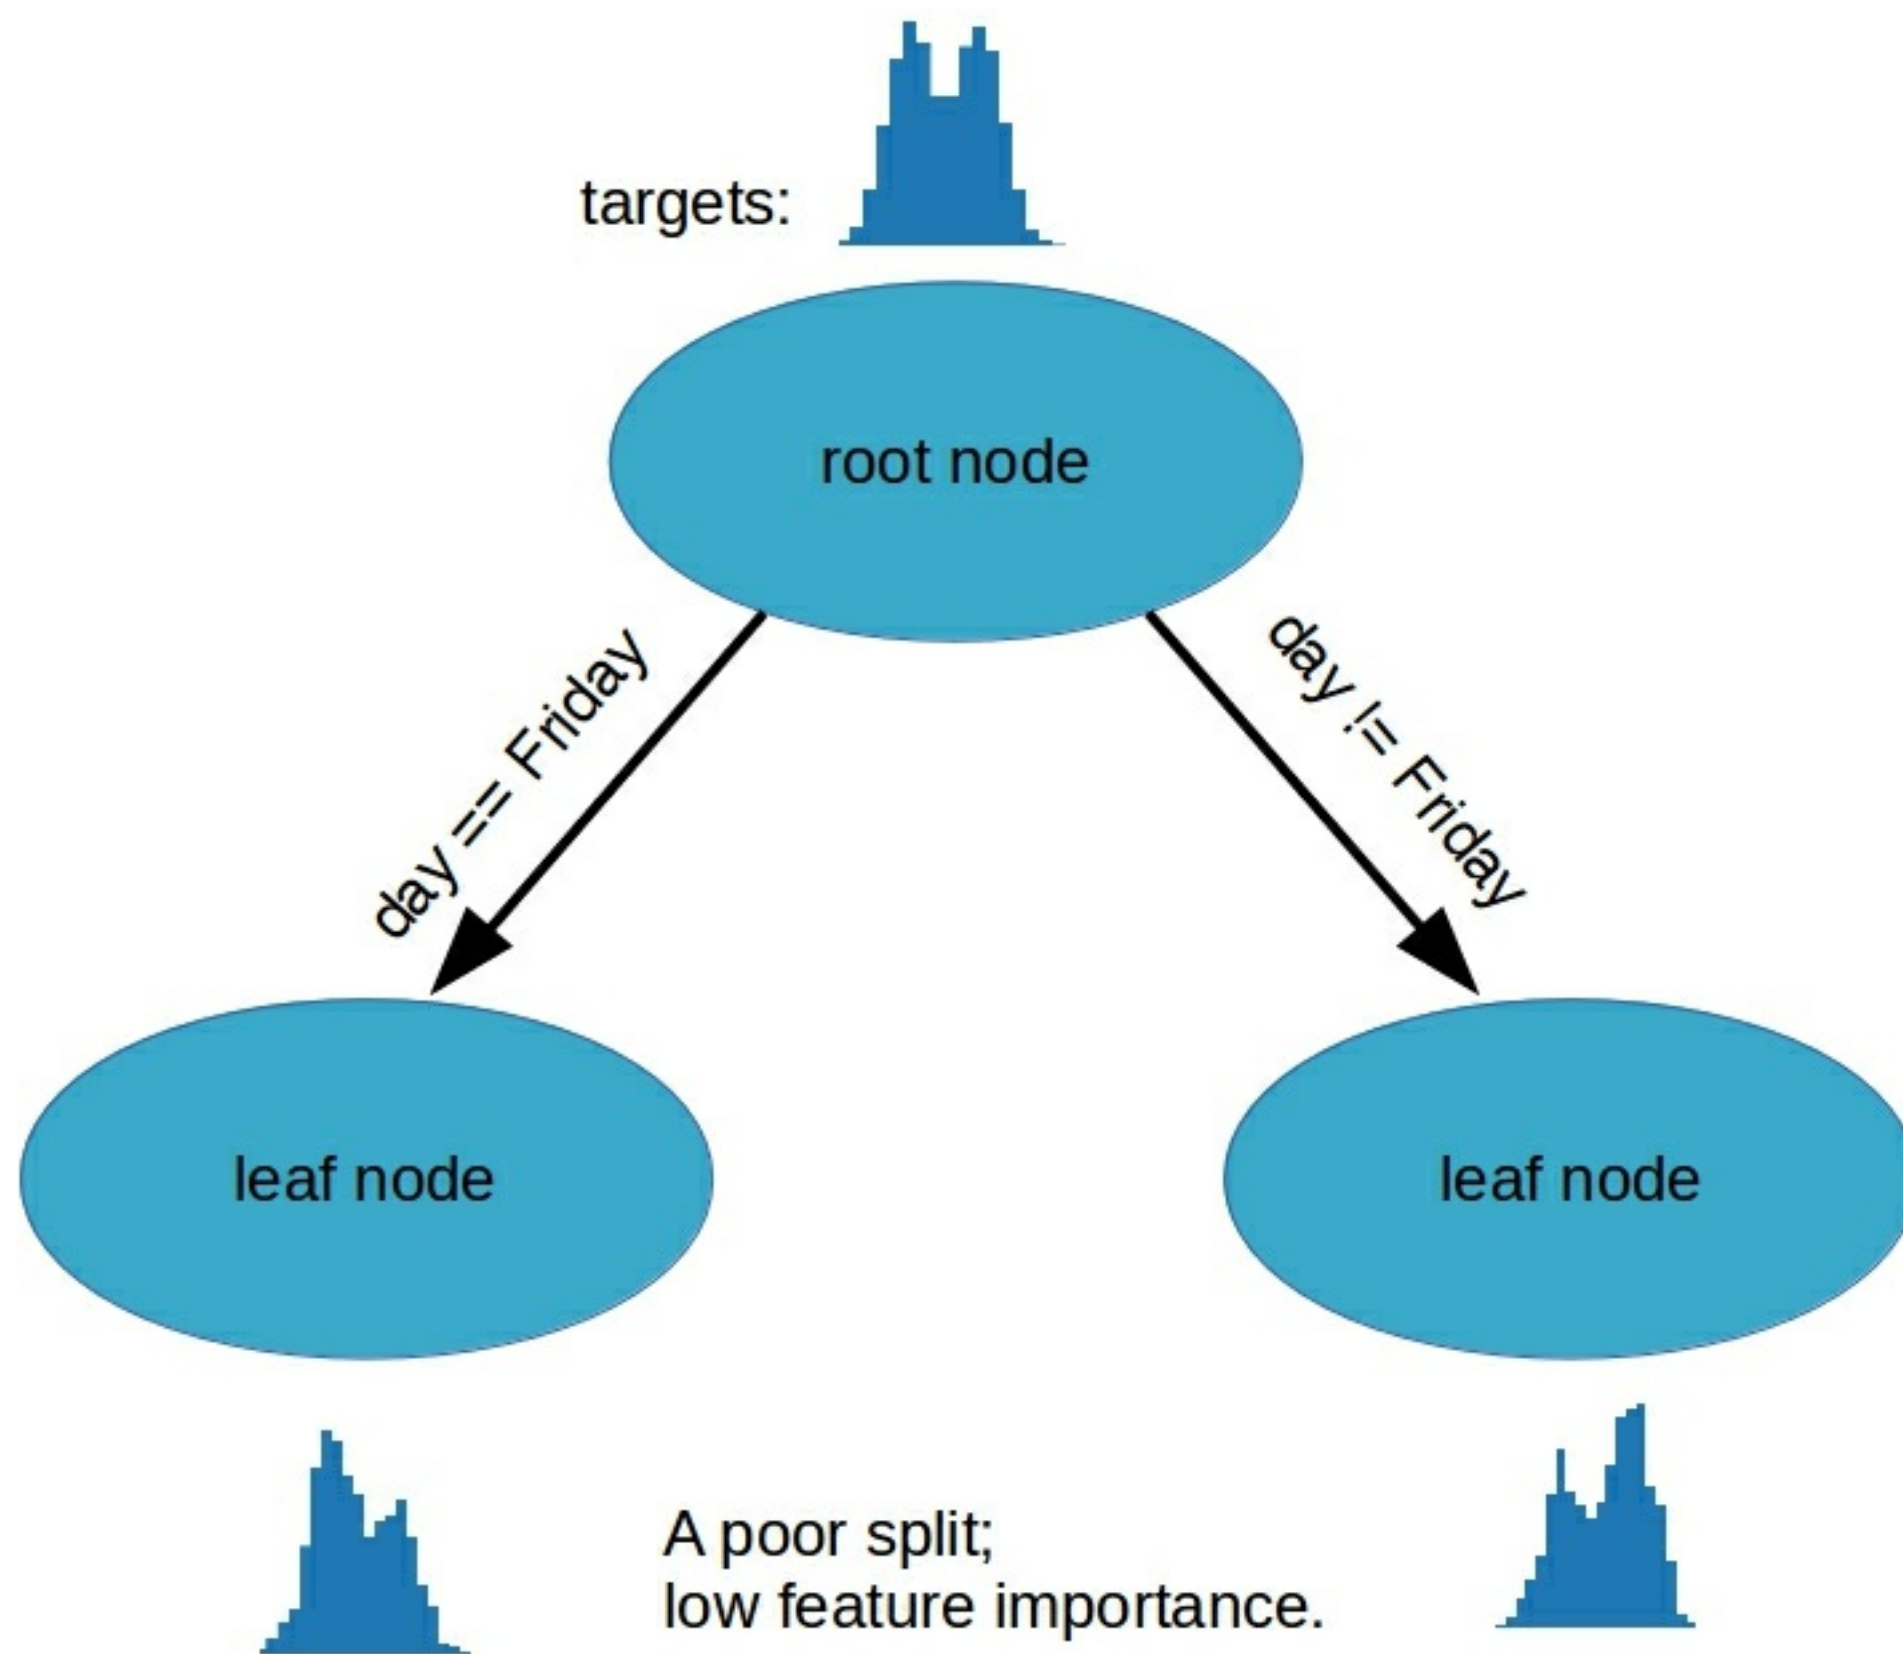
**Plant some random
forests!**

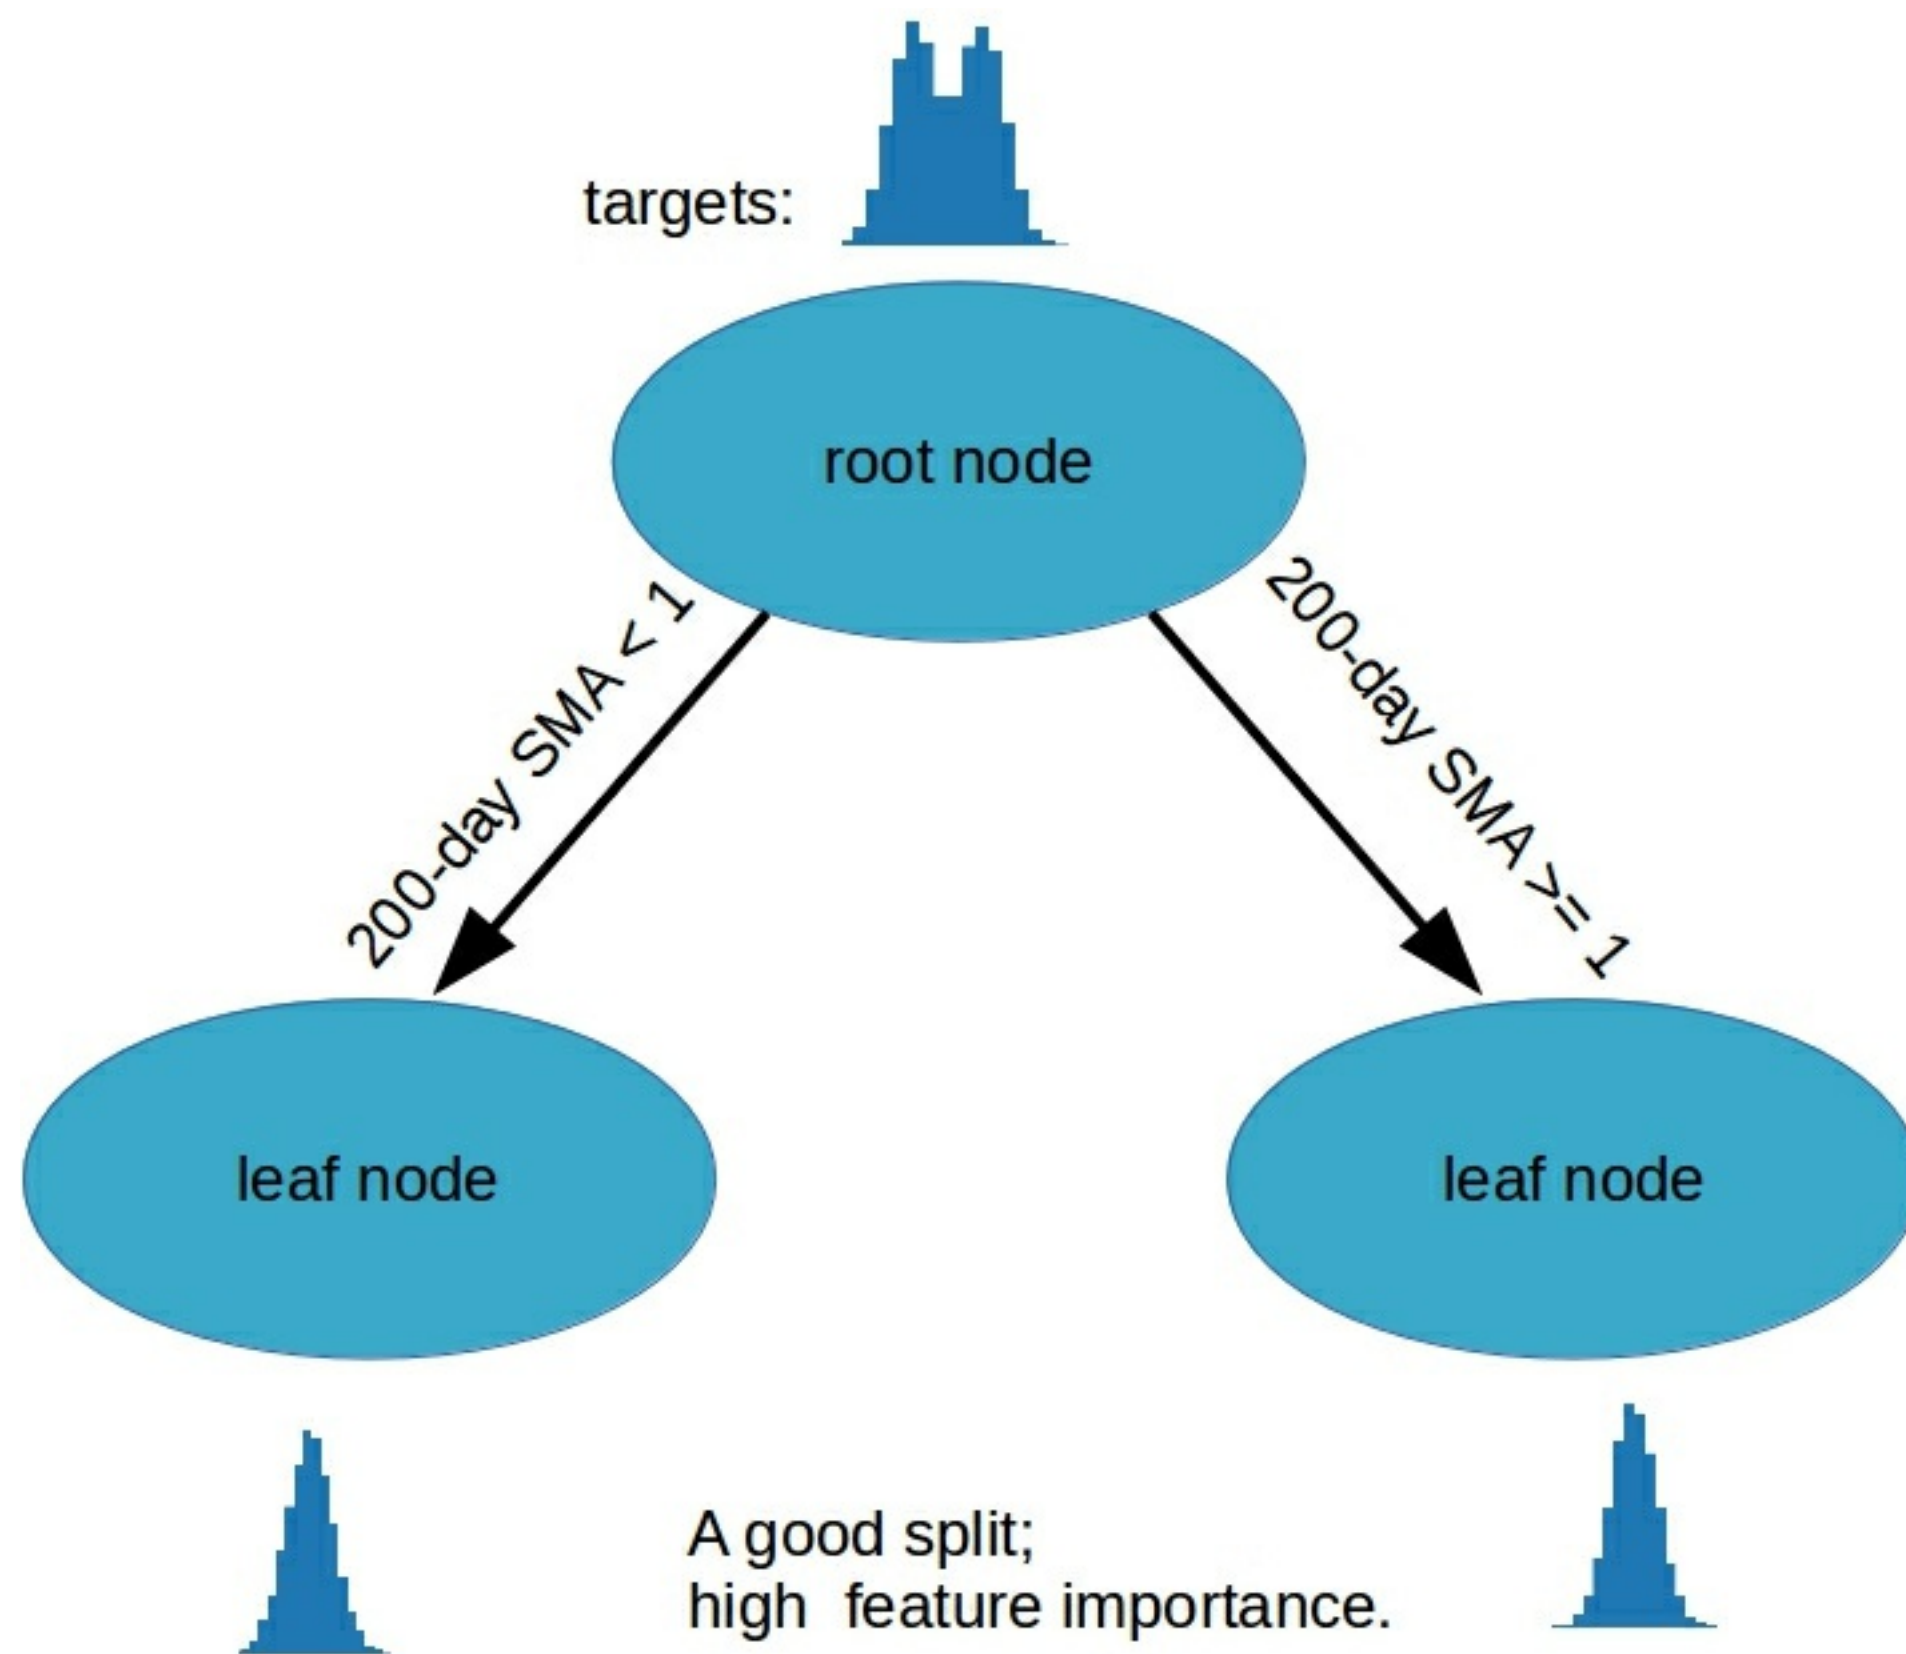


MACHINE LEARNING FOR FINANCE IN PYTHON

Feature importances and gradient boosting

Nathan George
Data Science Professor







Extracting feature importances

```
from sklearn.ensemble import RandomForestRegressor

random_forest = RandomForestRegressor()
random_forest.fit(train_features, train_targets)

feature_importances = random_forest.feature_importances_

print(feature_importances)
```

```
[0.07586547 0.10697602 0.12215955 0.23969227 0.29010304 0.0314028
 0.11977058 0.00276721 0.00246329 0.0026431  0.00615667]
```




Sorting and plotting


```
# feature importances from random forest model
importances = random_forest.feature_importances_

# index of greatest to least feature importances
sorted_index = np.argsort(importances)[::-1]

x = range(len(importances))
# create tick labels
labels = np.array(feature_names)[sorted_index]

plt.bar(x, importances[sorted_index], tick_label=labels)

# rotate tick labels to vertical
plt.xticks(rotation=90)
plt.show()
```




Linear models vs gradient boosting



Linear Regression

Gradient Boosting

<http://blog.kaggle.com/2017/01/23/a-kaggle-master-explains-gradient-boosting/>





Boosted models

Available boosted models:

- Gradient boosting
- Adaboost



Fitting a gradient boosting model

```
from sklearn.ensemble import GradientBoostingRegressor

gbr = GradientBoostingRegressor(max_features=4,
                                learning_rate=0.01,
                                n_estimators=200,
                                subsample=0.6,
                                random_state=42)

gbr.fit(train_features, train_targets)
```




MACHINE LEARNING FOR FINANCE IN PYTHON

Get boosted!